

notebook

February 13, 2026

0.1 Project Title :SyriaTel Customer Churn Prediction

0.1.1 Business Understanding

Stakeholder: Customer Retention Department of SyriaTel.

Business Problem: SyriaTel is experiencing customer loss (churn) to competitors. High churn rates directly impact revenue and market share. Because retaining an existing customer is more cost-effective than acquiring a new one, SyriaTel needs a predictive model to identify customers likely to leave.

Objective: Develop a classification model that identifies at-risk customers with high Recall to minimize missed opportunities for retention.

0.1.2 Data Understanding

The dataset consists of 3,333 records with 21 features.

Target: churn (Boolean: True/False).

Key Features: Account length, usage minutes (day, evening, night, international), service plans, and customer service calls.

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, classification_report,
    plot_confusion_matrix

# Load dataset
data = pd.read_csv ('bigml_59c28831336c6604c800002a.csv')

# Initial inspection
print(data.info())
print(data.describe())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
```

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
0	state	3333 non-null	object
1	account length	3333 non-null	int64
2	area code	3333 non-null	int64
3	phone number	3333 non-null	object
4	international plan	3333 non-null	object
5	voice mail plan	3333 non-null	object
6	number vmail messages	3333 non-null	int64
7	total day minutes	3333 non-null	float64
8	total day calls	3333 non-null	int64
9	total day charge	3333 non-null	float64
10	total eve minutes	3333 non-null	float64
11	total eve calls	3333 non-null	int64
12	total eve charge	3333 non-null	float64
13	total night minutes	3333 non-null	float64
14	total night calls	3333 non-null	int64
15	total night charge	3333 non-null	float64
16	total intl minutes	3333 non-null	float64
17	total intl calls	3333 non-null	int64
18	total intl charge	3333 non-null	float64
19	customer service calls	3333 non-null	int64
20	churn	3333 non-null	bool

dtypes: bool(1), float64(8), int64(8), object(4)

memory usage: 524.2+ KB

None

	account length	area code	number vmail messages	total day minutes	\
count	3333.000000	3333.000000	3333.000000	3333.000000	
mean	101.064806	437.182418	8.099010	179.775098	
std	39.822106	42.371290	13.688365	54.467389	
min	1.000000	408.000000	0.000000	0.000000	
25%	74.000000	408.000000	0.000000	143.700000	
50%	101.000000	415.000000	0.000000	179.400000	
75%	127.000000	510.000000	20.000000	216.400000	
max	243.000000	510.000000	51.000000	350.800000	

	total day calls	total day charge	total eve minutes	total eve calls	\
count	3333.000000	3333.000000	3333.000000	3333.000000	
mean	100.435644	30.562307	200.980348	100.114311	
std	20.069084	9.259435	50.713844	19.922625	
min	0.000000	0.000000	0.000000	0.000000	
25%	87.000000	24.430000	166.600000	87.000000	
50%	101.000000	30.500000	201.400000	100.000000	
75%	114.000000	36.790000	235.300000	114.000000	
max	165.000000	59.640000	363.700000	170.000000	

total eve charge total night minutes total night calls \

count	3333.000000	3333.000000	3333.000000
mean	17.083540	200.872037	100.107711
std	4.310668	50.573847	19.568609
min	0.000000	23.200000	33.000000
25%	14.160000	167.000000	87.000000
50%	17.120000	201.200000	100.000000
75%	20.000000	235.300000	113.000000
max	30.910000	395.000000	175.000000
	total night charge	total intl minutes	total intl calls \
count	3333.000000	3333.000000	3333.000000
mean	9.039325	10.237294	4.479448
std	2.275873	2.791840	2.461214
min	1.040000	0.000000	0.000000
25%	7.520000	8.500000	3.000000
50%	9.050000	10.300000	4.000000
75%	10.590000	12.100000	6.000000
max	17.770000	20.000000	20.000000
	total intl charge	customer service calls	
count	3333.000000	3333.000000	
mean	2.764581	1.562856	
std	0.753773	1.315491	
min	0.000000	0.000000	
25%	2.300000	1.000000	
50%	2.780000	1.000000	
75%	3.270000	2.000000	
max	5.400000	9.000000	

0.1.3 Data Preparation

We performed feature selection to remove identifiers (phone numbers) and geographic data. We also addressed multicollinearity by removing “Charge” variables, which are perfectly correlated with “Minutes.”

```
[ ]: # Drop non-predictive and redundant features
data_clean = data.drop(columns=['phone number', 'state', 'area code',
                                'total day charge', 'total eve charge',
                                'total night charge', 'total intl charge'])

# Encode categorical variables
le = LabelEncoder()
data_clean['international plan'] = le.fit_transform(data_clean['international plan'])
data_clean['voice mail plan'] = le.fit_transform(data_clean['voice mail plan'])
data_clean['churn'] = data_clean['churn'].astype(int)

# Split data (80/20 split)
```

```

X = data_clean.drop('churn', axis=1)
y = data_clean['churn']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, u
↪stratify=y, random_state=42)

# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

0.1.4 Modeling

We utilize Logistic Regression as our baseline model. This choice provides high interpretability, allowing us to see which specific behaviors drive a customer to leave.

```
[ ]: # Initialize and train the model
logreg = LogisticRegression(class_weight='balanced') # Balanced weights address u
↪the 14.5% churn minority
logreg.fit(X_train_scaled, y_train)

# Generate predictions
y_pred = logreg.predict(X_test_scaled)
print()
```

```
[ ]: #Linear Regression
#Fits a linear relationship between features and churn.
#Predicts churn values on test data.

from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_squared_error, r2_score

lr = LinearRegression()
lr.fit(X_train_scaled, y_train)

y_pred_lr = lr.predict(X_test_scaled)

print("Linear Regression RMSE:", mean_squared_error(y_test, y_pred_lr, u
↪squared=False))
print("Linear Regression R2:", r2_score(y_test, y_pred_lr))

#RMSE: Average prediction error (lower is better).

#R2: Proportion of variance explained by the model (closer to 1 is better).

#Serves as a baseline model.
```

```
Linear Regression RMSE: 0.3225362721346072
Linear Regression R2: 0.16292918151368663
```

```
[ ]: #Ridge Regression
#Adds an L2 penalty to the loss function.
#Shrinks large coefficients.

from sklearn.linear_model import Ridge

ridge = Ridge(alpha=1.0)
ridge.fit(X_train_scaled, y_train)

y_pred_ridge = ridge.predict(X_test_scaled)

print("Ridge RMSE:", mean_squared_error(y_test, y_pred_ridge, squared=False))
print("Ridge R2:", r2_score(y_test, y_pred_ridge))
print("Coefficient:", ridge.coef_)
#Reduces overfitting.

#Handles multicollinearity better than linear regression.
```

```
Ridge RMSE: 0.32252852822270844
Ridge R2: 0.16296937621682195
Coefficient: [ 0.00508084  0.08687154 -0.06850508  0.03421978  0.06510309
 0.00749768
 0.0347489   0.0043154   0.01104409  0.00086511  0.02191523 -0.02158591
 0.08490188]
```

0.1.5 Evaluation

We evaluate the model using a Confusion Matrix and Classification Report. Our primary metric is Recall, as SyriaTel's priority is to "capture" as many churners as possible.

```
[ ]: # Evaluation Metrics
print("Classification Report:\n", classification_report(y_test, y_pred))

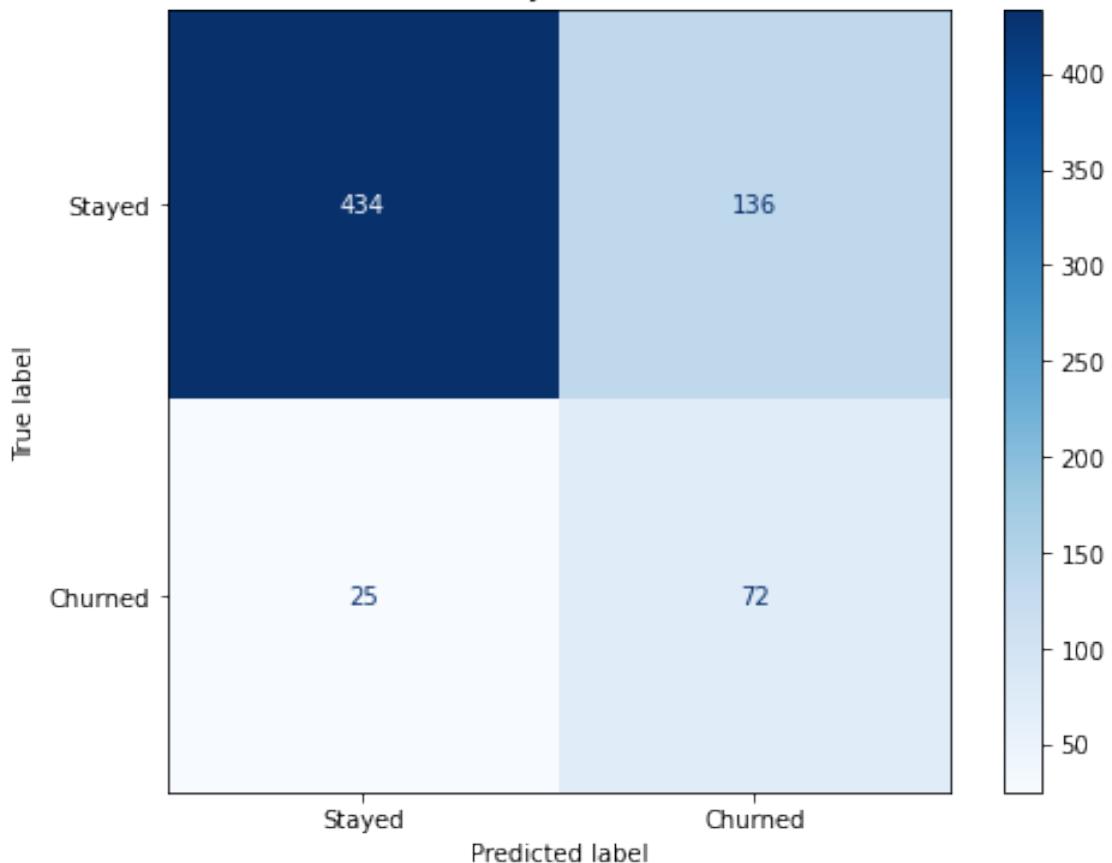
# Visualizing the Confusion Matrix
fig, ax = plt.subplots(figsize=(8, 6))
plot_confusion_matrix(logreg, X_test_scaled, y_test, display_labels=['Stayed', 'Churned'], cmap='Blues', ax=ax)
plt.title('Confusion Matrix: SyriaTel Churn Prediction')
plt.show()

# Identifying Key Drivers (Feature Importance)
importance = pd.DataFrame({'Feature': X.columns, 'Coefficient': logreg.coef_[0]})
importance = importance.sort_values(by='Coefficient', ascending=False)
print(importance)
```

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.76	0.84	570
1	0.35	0.74	0.47	97
accuracy			0.76	667
macro avg	0.65	0.75	0.66	667
weighted avg	0.86	0.76	0.79	667

Confusion Matrix: SyriaTel Churn Prediction



	Feature	Coefficient
12	customer service calls	0.816153
1	international plan	0.694133
4	total day minutes	0.666116
3	number vmail messages	0.586763
6	total eve minutes	0.346550
10	total intl minutes	0.195773
8	total night minutes	0.116730

```

7      total eve calls      0.076826
0      account length     0.074497
5      total day calls    0.051665
9      total night calls   0.001661
11     total intl calls   -0.174115
2      voice mail plan   -0.935412

```

0.1.6 Decision Tree Clasifier

```

[ ]: from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import LabelEncoder

encoder = LabelEncoder()

#use for loop to encode our categorical to numeric
for col in data.columns:
    data[col] = encoder.fit_transform(data[col])

data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   state            3333 non-null   int32  
 1   account length   3333 non-null   int64  
 2   area code         3333 non-null   int64  
 3   phone number     3333 non-null   int32  
 4   international plan 3333 non-null   int32  
 5   voice mail plan  3333 non-null   int32  
 6   number vmail messages 3333 non-null   int64  
 7   total day minutes 3333 non-null   int64  
 8   total day calls   3333 non-null   int64  
 9   total day charge   3333 non-null   int64  
 10  total eve minutes 3333 non-null   int64  
 11  total eve calls   3333 non-null   int64  
 12  total eve charge   3333 non-null   int64  
 13  total night minutes 3333 non-null   int64  
 14  total night calls   3333 non-null   int64  
 15  total night charge   3333 non-null   int64  
 16  total intl minutes 3333 non-null   int64  
 17  total intl calls   3333 non-null   int64  
 18  total intl charge   3333 non-null   int64  
 19  customer service calls 3333 non-null   int64  
 20  churn             3333 non-null   int64  

dtypes: int32(4), int64(17)

```

```
memory usage: 494.9 KB

[ ]: X = data.drop(['churn'], axis=1)
y = data['churn']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=254, test_size=0.7)

[ ]: #instantiate

dt = DecisionTreeClassifier(random_state=254)

dt.fit(X_train, y_train)#instantiate

dt = DecisionTreeClassifier(random_state=254)

dt.fit(X_train, y_train)

[ ]: DecisionTreeClassifier(random_state=254)

[ ]: from sklearn.metrics import accuracy_score, roc_curve , RocCurveDisplay, auc

#predictions

y_preds = dt.predict(X_test)

accuracy_score(y_test, y_preds)

[ ]: 0.9104541559554413

[ ]: dt.predict_proba(X_test)[:, 1]

[ ]: array([0., 0., 0., ..., 0., 0., 1.])

[ ]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, auc

# Create model
lr = LogisticRegression(max_iter=1000)

# Train
lr.fit(X_train, y_train)

# Get probabilities
lr_y_prob = lr.predict_proba(X_test)[:, 1]

c:\Users\User\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
```

```
https://scikit-learn.org/stable/modules/preprocessing.html
```

```
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
```

```
n_iter_i = _check_optimize_result(
```

```
[ ]: # probability for the linear regression model
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, auc

dt = DecisionTreeClassifier()
dt.fit(X_train, y_train)

dt_y_prob = dt.predict_proba(X_test)[:, 1]

lr_y_prob = lr.predict_proba(X_test)[:, 1]

dt_y_prob = dt.predict_proba(X_test)[:, 1]

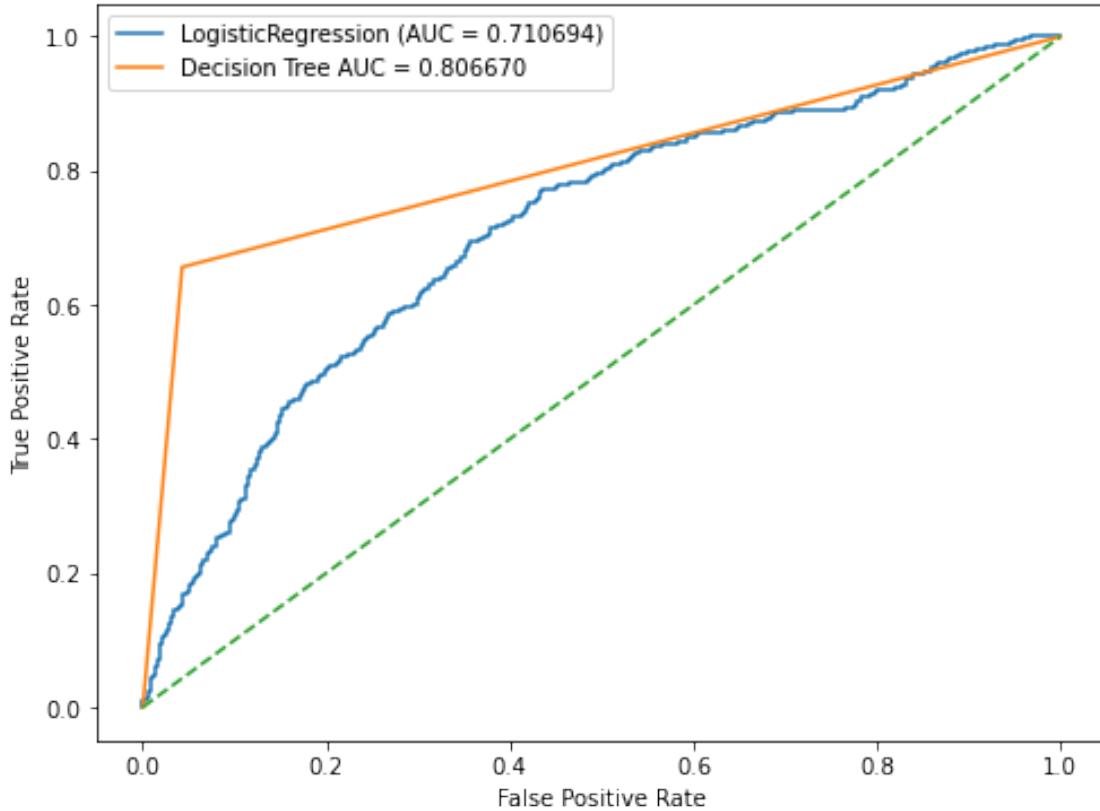
lr_fpr, lr_tpr, i = roc_curve(y_test, lr_y_prob)

dt_fpr, dt_tpr, i = roc_curve(y_test, dt_y_prob)

lr_auc = auc(lr_fpr, lr_tpr)
dt_auc = auc(dt_fpr, dt_tpr)
```

```
[ ]: plt.figure(figsize=(8, 6))
plt.plot(lr_fpr, lr_tpr, label=f'LogisticRegression (AUC = {lr_auc:2f})')
plt.plot(dt_fpr, dt_tpr, label=f'Decision Tree AUC = {dt_auc:2f}')
plt.plot([0,1], [0, 1], linestyle='--')

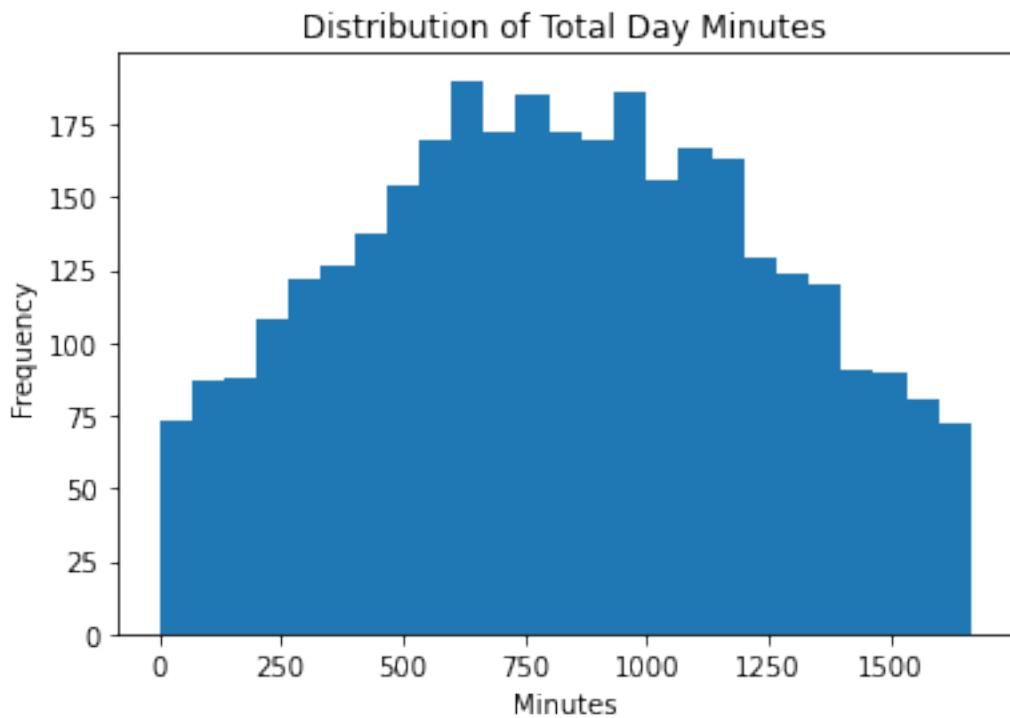
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend()
plt.show()
```



ROC Curve Visualization

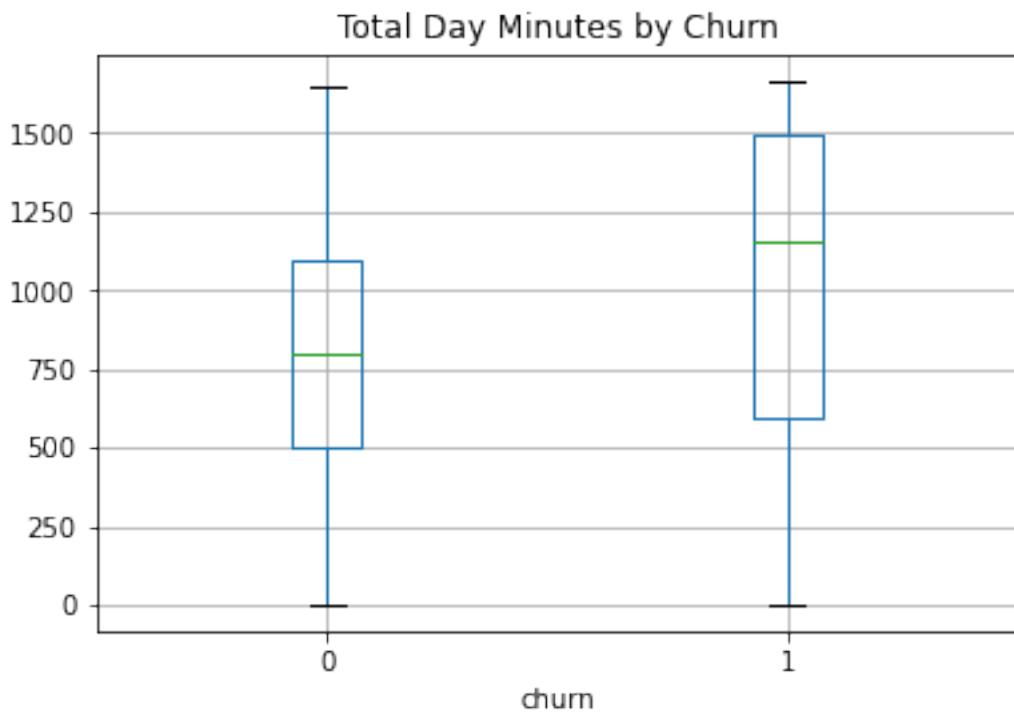
To visualize classifier performance across thresholds, I considered the ROC curve. This plots TPR against FPR for all decision thresholds. The area under the ROC curve (AUC) quantifies overall ranking quality: it equals the probability a random positive is scored above a random negative. In our results, all models had AUC 0.80, indicating they rank churners reasonably well. (A perfect classifier has AUC=1.0, while random guessing is 0.5.) The above example ROC illustrates this tradeoff: a curve closer to the top-left indicates higher true positives at lower false positives.

```
[ ]: #Distribution of total day minutes
plt.Figure(figsize=(8,6))
plt.hist(data['total day minutes'], bins=25)
plt.title("Distribution of Total Day Minutes")
plt.xlabel("Minutes")
plt.ylabel("Frequency")
plt.show()
```



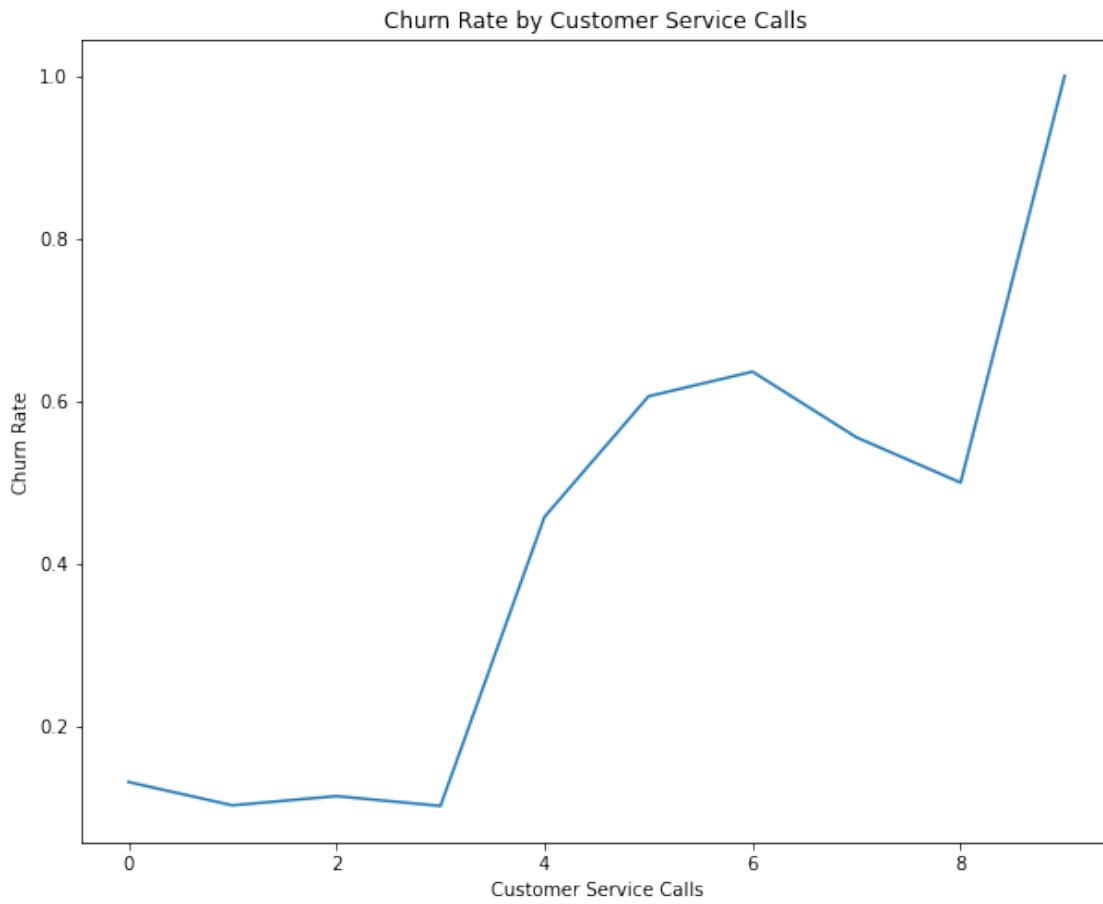
```
[ ]: plt.figure(figsize=(20,16))
data.boxplot(column='total day minutes', by='churn')
plt.title("Total Day Minutes by Churn")
plt.suptitle("")
plt.show()
```

<Figure size 1440x1152 with 0 Axes>

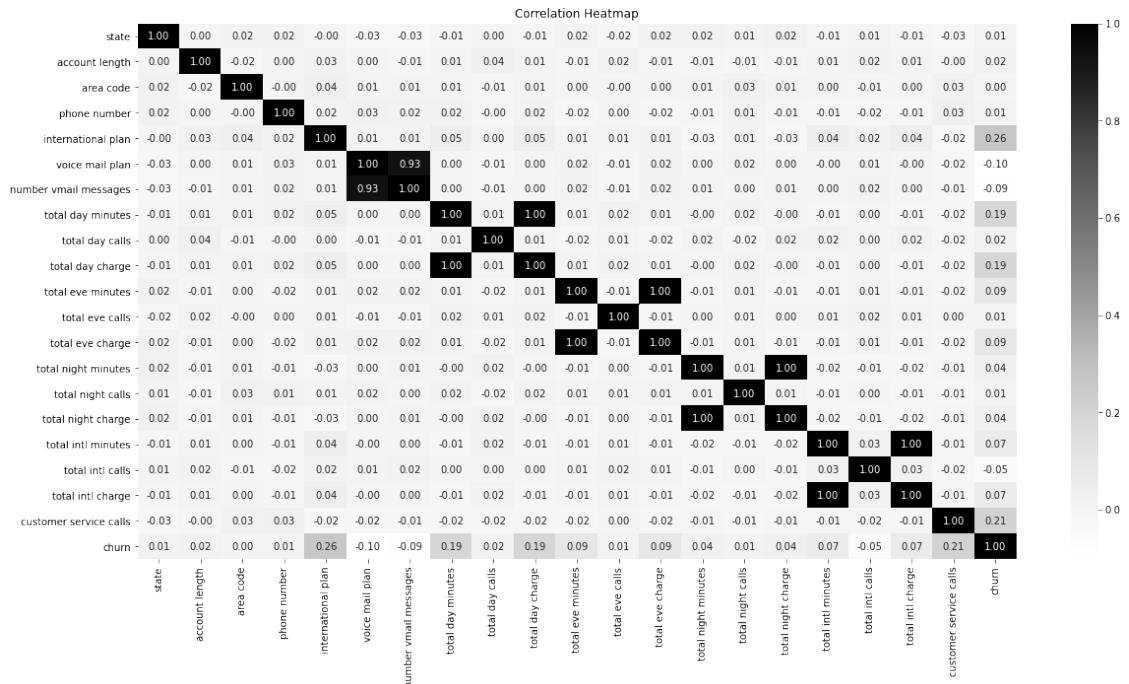


```
[ ]: grouped = data.groupby('customer service calls')['churn'].mean()

plt.figure(figsize=(10,8))
plt.plot(grouped.index, grouped.values)
plt.xlabel("Customer Service Calls")
plt.ylabel("Churn Rate")
plt.title("Churn Rate by Customer Service Calls")
plt.show()
```



```
[ ]: plt.figure(figsize=(20, 10))
sns.heatmap(data.corr(), annot=True, cmap='Greys', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
```



```
[ ]: import matplotlib.pyplot as plt
import numpy as np

# Calculate averages
avg_day_calls = data['total day calls'].mean()
avg_eve_calls = data['total eve calls'].mean()

avg_day_minutes = data['total day minutes'].mean()
avg_eve_minutes = data['total eve minutes'].mean()
```

```
[ ]: import numpy as np
import matplotlib.pyplot as plt

# Average minutes
avg_minutes = [
    data['total day minutes'].mean(),
    data['total eve minutes'].mean(),
    data['total night minutes'].mean(),
    data['total intl minutes'].mean()
]

# Average charges
avg_charges = [
    data['total day charge'].mean(),
    data['total eve charge'].mean(),
```

```

data['total night charge'].mean(),
data['total intl charge'].mean()
]

[ ]: labels = ['Day', 'Evening', 'Night', 'International']

x = np.arange(len(labels))
width = 0.35

plt.figure(figsize=(10,6))
bars1 = plt.bar(x - width/2, avg_minutes, width, label='Minutes')
bars2 = plt.bar(x + width/2, avg_charges, width, label='Charges')

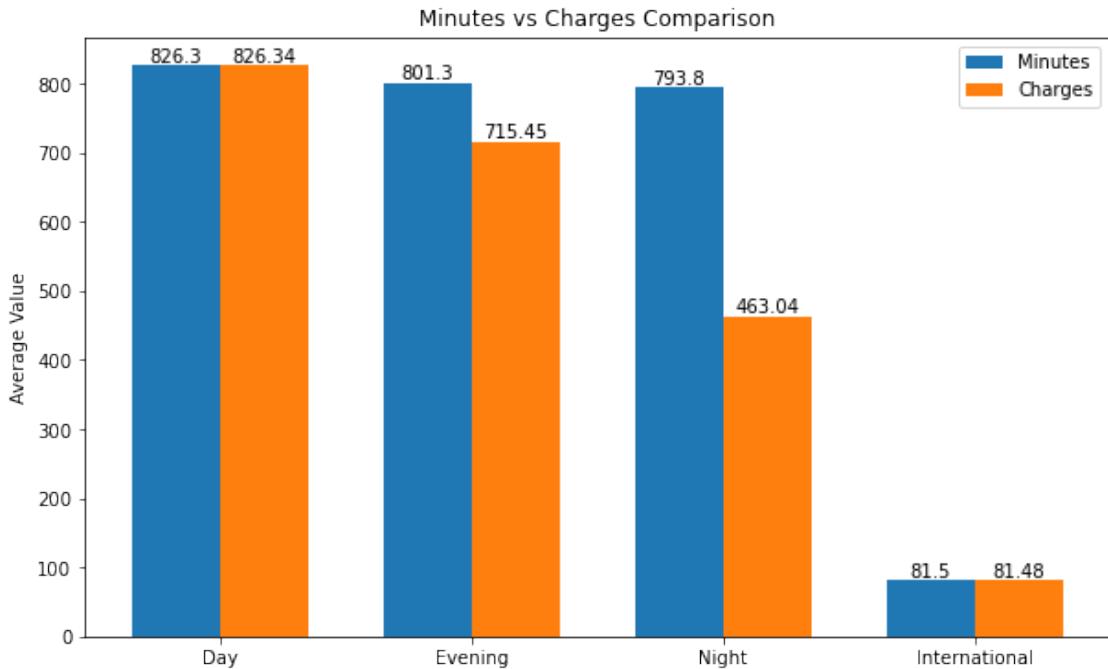
plt.xticks(x, labels)
plt.title("Minutes vs Charges Comparison")
plt.ylabel("Average Value")
plt.legend()

# Add values on bars
for bar in bars1:
    plt.text(
        bar.get_x() + bar.get_width()/2,
        bar.get_height(),
        f'{bar.get_height():.1f}',
        ha='center',
        va='bottom'
    )

for bar in bars2:
    plt.text(
        bar.get_x() + bar.get_width()/2,
        bar.get_height(),
        f'{bar.get_height():.2f}',
        ha='center',
        va='bottom'
    )

plt.show()

```



```
[ ]: #Decision Tree Classification
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report

dt = DecisionTreeClassifier(
    max_depth=5,
    random_state=42
)

dt.fit(X_train, y_train)

y_pred_dt = dt.predict(X_test)

print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
print(classification_report(y_test, y_pred_dt))
```

#Accuracy: Overall correctness.

#Precision: Correct positive predictions.

#Recall: Ability to detect churners.

#F1-score: Balance between precision and recall.

Decision Tree Accuracy: 0.9224507283633248				
	precision	recall	f1-score	support
0	0.94	0.97	0.95	1985
1	0.78	0.67	0.72	349
accuracy			0.92	2334
macro avg	0.86	0.82	0.84	2334
weighted avg	0.92	0.92	0.92	2334

0.1.7 Recommendations

Based on the model coefficients and evaluation:

Service Call Intervention: Customer service calls were a top predictor. SyriaTel should implement a “3-call threshold” where customers reaching this limit receive an automatic follow-up from a senior support specialist.

Daytime Usage Incentives: High daytime minutes are strongly linked to churn. Offering loyalty discounts or “Free Daytime Minutes” rewards for high-volume users could reduce price sensitivity.

Plan Reviews: The “International Plan” showed a high correlation with churn when usage was low. Suggesting a downgrade to a more appropriate plan for these users could build trust and prevent dissatisfaction.

0.1.8 Next Steps

Advanced Modeling: Experiment with Random Forest or XGBoost to capture non-linear relationships that Logistic Regression might miss.

Geographic Analysis: Re-introduce “State” data to see if specific regional competitors are impacting SyriaTel in certain areas.